

Parallelizing Programs

- Goal: speed up programs using multiple processors/cores

(Sequential) Matrix Multiplication

```
double A[n][n], B[n][n], C[n][n] // assume n x n
for i = 0 to n-1
    for j = 0 to n-1
        double sum = 0.0
        for k = 0 to n-1
            sum += A[i][k] * B[k][j]
        C[i][j] = sum
```

Question: how can this program be parallelized?

Steps to parallelization

- First: find parallelism
 - Concerned about what can *legally* execute in parallel
 - At this stage, expose as much parallelism as possible
 - Partitioning can be based on data structures or by function

Note: other steps are architecture dependent

Finding Parallelism in Matrix Multiplication

- Can we parallelize the inner loop?

Finding Parallelism in Matrix Multiplication

- Can we parallelize the inner loop?
 - No, because *sum* would be written concurrently

Finding Parallelism in Matrix Multiplication

- Can we parallelize the inner loop?
 - No, because *sum* would be written concurrently
- Can we parallelize the outer loops?

Finding Parallelism in Matrix Multiplication

- Can we parallelize the inner loop?
 - No, because *sum* would be written concurrently
- Can we parallelize the outer loops?
 - Yes, because the read and write sets are independent for each iteration (i,j)
 - Read set for process (i,j) is *sum*, $A[i][k=0:n-1]$, $B[k=0:n-1][j]$
 - Write set for process (i,j) is *sum*, $C[i][j]$
 - Note: we have the option to parallelize just one of these loops

Terminology

- *co* statement: creates parallelism

co $i := 0$ to $n-1$

Body

oc

- Meaning: n instances of body are created and executed concurrently until the end of the *co* (i.e., at the *oc*)
- Implementation: fork n threads, join them at the *oc*

Need to understand what processes/threads are!

Processes

- History: OS had to coordinate many activities
 - Example: deal with multiple users (each running multiple programs), incoming network data, I/O interrupts
- Solution: Define a model that makes complexity easier to manage
 - Process (thread) model

What's a process?

- Informally: program in execution
- Process encapsulates a physical processor
 - everything needed to run a program
 - code (“text”)
 - registers (PC, SP, general purpose)
 - stack
 - data (global variables or dynamically allocated)
 - files
- NOTE: a process is sequential

Examples of Processes

- Shell: creates a process to execute command

lectura:> ls foo

(shell creates process that executes “ls”)

lectura:> ls foo & cat bar & more

(shell creates three processes, one per command)


- OS: creates a process to manage printer
 - process executes code such as:
 - wait for data to come into system buffer
 - move data to printer buffer

Creating a Process


- Must somehow specify code, data, files, stack, registers
- Ex: UNIX
 - Use the `fork()` system call to create a process
 - Makes an **exact** duplicate of the current process
 - (returns 0 to indicate child process)
 - Typically `exec()` is run on the child

We will not be doing this (systems programming)


Example of Three Processes



Process 0



Process 1



Process 2

OS switches between the three processes (“multiprogramming”)

Review: Run-time Stack

```
A(int x) {  
    int y = x;  
    if (x == 0) return;  
    else return A(y-1) + 1;  
}
```

```
B() {  
    int z;  
    A(1);  
}
```

y (0)
x (0)
y (1)
x (1)
z

Decomposing a Process

- Process: everything needed to run a program
- Consists of:
 - Thread(s)
 - Address space


Thread

- Sequential stream of execution
- More concretely:
 - program counter (PC)
 - register set
 - stack
- Sometimes called lightweight process


Address Space

- Consists of:
 - code
 - data
 - open files
- Address space can have > 1 thread
 - threads share code, data, files
 - threads have separate stacks, register set

One Thread, One Address Space



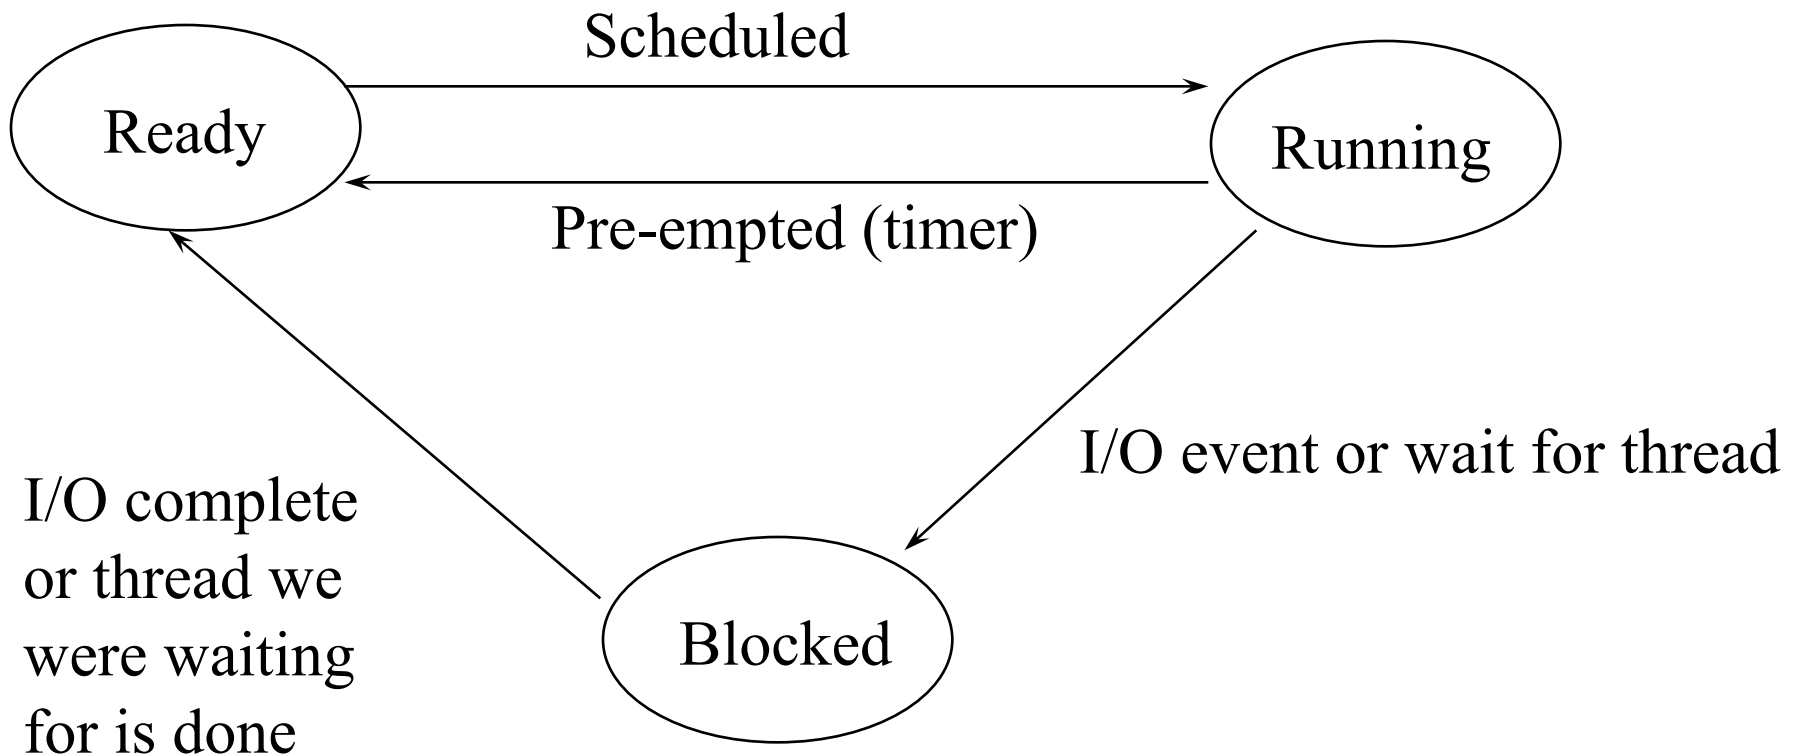
Many Threads, One Address Space



Thread States

- Ready
 - eligible to run, but another thread is running
- Running
 - using CPU
- Blocked
 - waiting for something to happen

Thread State Graph



Scheduler

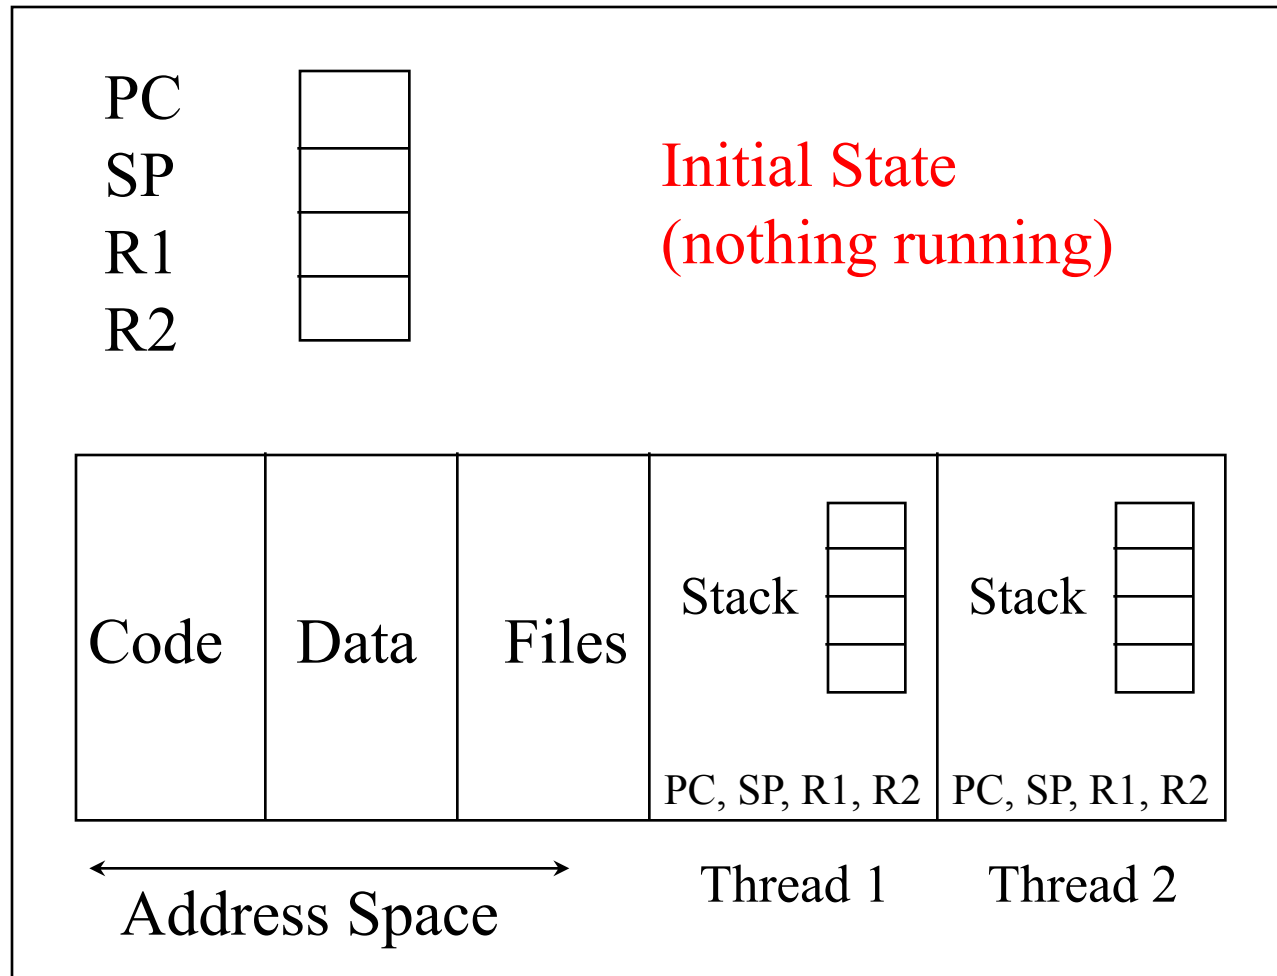
- Decides which thread to run
 - (from ready list only)
- Chooses from some algorithm
- From our point of view, the scheduler is something we cannot control
 - We have no idea which thread will be run, and our programs must not depend on execution order of two ready threads

Context Switching

- Switching between 2 threads
 - change PC to current instruction of new thread
 - **might need to restart old thread in the future**
 - must save exact state of first thread
- What must be saved?
 - registers (including PC and SP)
 - what about stack itself?


Multiple Threads, One Machine (Single Core)

Machine




Multiple Threads, One Machine (Single Core)

Machine




Multiple Threads, One Machine (Single Core)

Machine



Multiple Threads, One Machine (Single Core)

Machine



Why Save Registers?

- code for Thread 0

foo()

$x := x + 1$

$x := x * 2$

Assembly code:

$R1 := R1 + 1$ /* !! */

$R1 := R1 * 2$

- code for Thread 1

bar()

$y := y + 2$

$y := y - 3$

Assembly code:

$R1 := R1 + 2$

$R1 := R1 - 3$

Suppose context switch
occurs after line “!!”

Matrix Multiplication, n^2 threads

```
double A[n][n], B[n][n], C[n][n] // assume n x n
```

```
co i = 0 to n-1 {
```

```
    co j = 0 to n-1 {
```

```
        double sum = 0.0
```

```
        for k = 0 to n-1
```

```
            sum += A[i][k] * B[k][j]
```

```
        C[i][j] = sum
```

```
    }
```

```
}
```

We already argued the two outer
“for” loops were parallelizable

Steps to parallelization

- Second: control granularity
 - Must trade off advantages/disadvantages of fine-granularity
 - Advantages: better load balancing, better scalability
 - Disadvantages: process/thread overhead and communication
 - Combine small processes into larger ones to coarsen granularity
 - Try to keep the load balanced

Matrix Multiplication, n threads

```
double A[n][n], B[n][n], C[n][n] // assume n x n
```

```
co i = 0 to n-1 {
```

```
  for j = 0 to n-1 {
```

```
    double sum = 0.0
```

```
    for k = 0 to n-1
```

```
      sum += A[i][k] * B[k][j]
```

```
    C[i][j] = sum
```

```
  }
```

```
}
```

This is plenty of parallelization
if the number of cores is $\leq n$

Matrix Multiplication, p threads

```
double A[n][n], B[n][n], C[n][n] // assume n x n
```

```
co t = 0 to p-1 {
```

```
    startrow = t * n / p; endrow = (t+1) * n/p - 1
```

```
    for i = startrow to endrow
```

```
        for j = 0 to n-1 {
```

```
            double sum = 0.0
```

```
            for k = 0 to n-1
```

```
                sum += A[i][k] * B[k][j]
```

```
            C[i][j] = sum
```

```
        }
```

```
    }
```


Steps to parallelization

- Third: distribute computation and data
 - Assign which processor does which computation
 - The `co` statement does *not* do this
 - If memory is distributed, decide which processor stores which data (why is this?)
 - Data can be replicated also
 - Goals: minimize communication and balance the computational workload
 - Often conflicting goals

Steps to parallelization

- Fourth: synchronize and/or communicate
 - If shared-memory machine, synchronize
 - Both mutual exclusion and sequence control
 - Locks, semaphores, condition variables, barriers, reductions
 - If distributed-memory machine, communicate
 - Message passing
 - Usually communication involves implicit synchronization

Parallel Matrix Multiplication---

Distributed-Memory Version

```
process worker [i = 0 to p-1] {  
  double A[n][n], B[n][n], C[n][n] // wasting space!  
  startrow = i * n / p; endrow = (i+1) * n/p - 1  
  if (i == 0) {  
    for j = 1 to p-1 {  
      sr= j * n / p; er = (j+1) * n/p - 1  
      send A[sr:er][0:n-1], B[0:n-1][0:n-1] to process j  
    }  
  else  
    receive A[startrow:endrow][0:n-1], B[0:n-1][0:n-1] from 0
```

Parallel Matrix Multiplication---

Distributed-Memory Version

```
for i = startrow to endrow
```

```
  for j = 0 to n-1 {
```

```
    double sum = 0.0
```

```
    for k = 0 to n-1
```

```
      sum += A[i][k] * B[k][j]
```

```
    C[i][j] = sum
```

```
  }
```

```
// here, need to send my piece back to master
```

```
// how do we do this?
```

```
} // end of process statement
```

Adaptive Quadrature: Sequential Program

```
double f() {  
    ....  
}  
double area(a, b)  
    c := (a+b)/2  
    compute area of each half and area of whole  
    if (close)  
        return area of whole  
    else  
        return area(a,c) + area (c,b)
```

Adaptive Quadrature: Recursive Program

```
double f() {  
    ....  
}  
double area(a, b)  
    c := (a+b)/2  
    compute area of each half and area of whole  
    if (close)  
        return area of whole  
    else  
        co leftArea = area(a,c) // rightArea = area (c,b) oc  
        return leftArea + rightArea
```

Challenge with Adaptive Quadrature

- For efficiency, must control granularity (step 2)
 - Without such control, granularity will be too fine
 - Can stop thread creation after “enough” threads created
 - Hard in general, as do not want cores idle either
 - Thread implementation can perform work stealing
 - Idle cores take a thread and execute that thread, but care must be taken to avoid synchronization problems and/or efficiency problems

Steps to parallelization

- Fifth: assign processors to tasks (only if using task and data parallelism)
 - Must also know dependencies between tasks
 - Usually task parallelism used if limits of data parallelism are reached

Steps to parallelization

- Sixth: parallelism-specific optimizations
 - Examples: message aggregation, overlapping communication with computation

Steps to parallelization

- Seventh: acceleration
 - Find parts of code that can run on GPU/Xeon Phi/etc., and optimize those parts
 - Difficult and time consuming
 - But may be quite worth it